

# Machine Learning A-Z

## Real World Data Science Challenge Steps:

- Formulate the problem
- What choice of evaluation
- What data to be used
- Build models
- Evaluate the model

## Learning Models

*There is no free lunch. Each model expose its assumption on data and reduces generalizability. Most widely used are XGBoost and NN.*

**Linear models:** separates with linear lines

- Pro: Good with sparse high dimensional data
- Con: Too simple and might introduce bias
- Ex: SVM, linear regression, logistic regression
- Lib: Sikitlearn, Vowpal
- Consideration: requires feature scaling, outlier removing

**Tree based models:** divide and conquer approach

- Pro: flexibility in assembling, very good starting points, no need for scaling
- Con: don't do well with linear data
- Ex: Decision trees, random forest, gradient boosting trees
- Lib: Sikitlearn, XGBoost, LightGBM

**Nearest neighbors:** looking at neighbors for insight

- Pro: can be used for feature generation, fast
- Con: don't do well with linear data
- Ex: K-NN, K-means clustering
- Lib: Sikitlearn
- Consideration: requires feature scaling, outlier removing

**Neural networks:** looking at neighbors for insight

- Pro: smooth boundaries, black box, good for image, sound, text, sequence
- Con: don't do well with linear data
- Ex: convolutional neural network,
- Lib: Tensorflow, keras, pytorch, lasgna
- Consideration: requires feature scaling, outlier removing

**Random Forests:** an ensembler for decision trees

Randomized based on the number of columns and rows included

- Pro: bias stays at the level of DT, while variance decreases, quick and dirty solution, good for feature selection
- Con: features are not interpretable, not good for smaller data sets. In regression, range is limited to available data.
- Lib: Sikitlearn

## Feature Processing

Depending on the type of feature, there are some well-known techniques. Choice of model is also influential.

### Numerical data:

**Processing:**

- **Scaling:** sklearn.preprocessing.StandardScaler
- **Clipping:** numpy.clip (for outliers)
- **Rank transformation:** scipy.stat.rankdata (for outliers)
- **Transformation:** log, sqrt (shrinking outliers)

**Engineering:**

- **Interaction:** \*//+ between features (dig data)
- **Decimal points:** 4.99 → 0.99
- **Group statistics:** std, mean, max, min of the group this observation is in (use groupby, agg).

### Categorical data:

**Processing:**

- **Label Encoding:** generate labels from 0 to n-1 level, good for ordinals and linear models. Sklearn.preprocessing.labelencoder
- **Factorization:** for non-ordinals, pandas.factorize
- **Frequency encoding:** giving the ratio of the number of that class as the value for it, when frequency might be correlated with output. Use rank when some frequencies are too close.
- **OneHotCoder:** generates scaled dummies, good for linear classifiers, and non linear with target, sklearn.preprocessing.onehotencoder, you can also use pd.get\_dummies(df, drop\_first=True)

**Engineering:**

- **Interactions:** .str+.str
- **Mean encoding:** each class has a value equal to mean of target in that class multiplied by its encoded label.

### Datetime data:

**Engineering:**

- **Time moments and periods:** year, month, ...
- **Time past since or until:** last holiday, last time bought
- **Time delta between two datetime columns,**

### Text data:

**Processing:**

- **lowercasing:** .lower()
- **Lemmatization/stemming:** getting the roots of words
- **Stopword cleaning:** taking repetitive words out

**Engineering:**

# Machine Learning A-Z

- **Bag of words:** one column for each word and see its repetition in each row, `sklearn.feature_extraction.text.CountVectorizer`
- **TFIDF**, normalizing words bag of words, both row-wise and column wise (how important is feature in that sentence, and how repetitive is it in all data) `Sklearn.feature_extraction.TfidfVectorizer`
- **N-gram:** bag of words for N neighbor words, `CountVectorizer` has “n-range”, use it for vectorizing.
- **Word2vec:** transfers words to a complicated space. Meaning of vectors are not known, but space is much smaller than bag of words. Similar vectors have similar meaning.

## Image data:

### Engineering:

- **CNN:** similar to word2vec but for images, you can either train from scratch or use another network with similar object and **finetune** the model.

## Coordinate data:

### Engineering:

- **Distance to an important point, old building,**

## Handling missing values:

Look into histogram, any unusual bump might be missing values encoded.

### Imputations:

- A value far out of the min and max range, (good for tree based)
- Mean, median (good for non-trees)
- A categorical value (`is_missing`)

## Feature Selection

- **Unsupervised learning:** Use PCA, NMF, SVD to choose feature space with more info,

- **Feature importance:** use Random Forest, XGBoost, to get the feature importance

## EDA and Data Cleaning

Helps understanding the data and getting it ready for deploying to classifiers.

### EDA Tips:

#### Higher level understanding:

- `df.dtypes()`, `df.info()`, `df.describe()`, `df.isnull()`, `df.head()`, `df.tail()`, `select_dtypes(include=[int])`

#### Visualization

- single variable: `plt.hist()`, `plt.plot(df, '.')` (good if standard scaled, to see how data changes through all rows)
- multiple variables: `plt.scatter(f1, f2, c=target)`, `plt.matshow(df.corr())`, `df.mean().sort_values().plot(style='')`

### Cleaning Tips:

- dropping rows with similar value in all rows: `df=[~df.nunique(axis=1)==1]`
- `df.T.drop_duplicates()`
- for categorical duplicates with different names: for `f` in `categorical_feat`: `dff=df[f].factorize()` then the above comment
- `select_dtypes(include=[int])`

## Validation and overfitting

Does model generalize well?

### Validation:

#### Types:

- Holdout: dividing into train test
- K-fold: repeated hold-out then average
- Leave one out: only when model is very fast, you can have `k=len(data)`

Make validation stratified if the classes are skewed and there are many features in one vs other

### Data split strategies:

Validation should mimic the train test split (this is important for kaggle). But for real-world, validation should mimic real world. Types of validation:

- Random split
- Time-wise split
- By id

### Data leaks:

(related to kaggle) Unexpected info in data that allows to estimate the target without building a model. Types of data leaks:

- Time series: future picking,
- Meta data: e.g. time and place pictures taken

### Metrics:

#### Metric vs Loss:

Metric is how we assess the quality of technique, loss is what model tries to minimize. Sometimes they overlap (like regression MSE).

### Validation (evaluation) metrics:

We need to estimate how we work on real-world situations. Different metrics results in different hyperplanes.

#### 1- Regression Metrics:

- Mean square error (**MSE**). Mean value minimizes it, default for most regression,
- Root mean square error (`sqrt(MSE)`)
- **R-squared**, how much better than a baseline (`y=mean`) MSE prediction we are?
- Mean absolute error (**MAE**), where errors are linearly penalized. Median minimizes it.

# Machine Learning A-Z

## XGBoost, LightGBM

- Maxdepth (hoe many level), subsample (bagging fraction), colsample\_by\_level, colsample\_by\_tree, eta (learning rate ), num\_round (how many iteration)
- Min\_child\_weight (min data in leaf) most important parameter

## RandomForest

Random forest trees are independent as opposed to XGBoost.

- N\_estimator : number of trees, makes model better and saturate in a place.
- Max\_depth, how much deep you can go,
- Max\_featue: number of features used in each tree
- Min\_sample\_leaf
- Criterion: gene works good
- N\_jobs: number of cores on your machine.

## Neural Networks

- Number of neurons per layer, number of layers, batch size,
- Optimizer: either choose gradient descent + momentum, or adam
- Learning rate,
- Neural net regularizes:
  - L2/L1 for weights
  - Dropout, drop connections (do this close to the last layer)
  - Statistic drop connect

## SVM

- Regularization, regularization type.

## Ensembling

The combination performed on the prediction of different models on different part of data set on your validation set. There are different types:

- **Averaging:** averaging the predictions,
- **Weighted averaging:** averaging predictions by weighting models with higher quality,
- **Conditional:** take prediction for a range of target from one model, and the rest from other models
- **Bagging:** averaging slightly different independent version of the same model (random forest does this)
- **Boosting:** each model is built to explain the unexplained part of target for previous model(s). Gradient boosted trees do this.
- **Stacking:** training a model on the predicted value vs actual value of the validation target. Model is usually simpler. Called meta modelling. We divide data 50/50 between train/val.

### Tips on ensembling:

- First ensemble several gradient boosted trees,
- 2-3 neural network,
- A few random forests,
- One linear regression
- A couple of K-NNs
- SVM with non-linear kernel
- Next layers have simpler algorithms,
- In next layers, feature engineering can be difference between predictions ( $y_{model1} - y_{model2}$ ).

- Mean square percentile error(**MSPE**), relative error minimization, depending on the size of target (error/target is minimized), weighted mean of target value minimizes it,
- Mean absolute percentage error(**MAPE**), weighted median of target value minimizes it.
- Root mean square logarimic error, weighted mean in logspace minimizes it.

### 2- Classification metrics:

- **Accuracy:** fraction of correctly predicted, non-immune to skewed classes, non-caring about level of confidence in the prediction
- **Logloss:** requires soft prediction (probability of target). Penalize when confident and wrong. Class ratio (class A/total) minimizes it.
- **Are under curve:** tries all possible threshold and find the accuracy and aggregate on it. Its base is 50%.
- **Cohen's Kappa:** good for skewed classes,  $1 - (1 - \text{accuracy}) / (1 - \text{baseline})$ . You need to combine weighted error matrix and confusion matrix.

### Metric optimization:

- Models optimize loss, not metrics. MSE and Logloss can be both metric and loss. But the rest needs optimization.

### Early Stopping, golden method for optimization:

- We keep training the model and calculate the cross validation metrics, until the metric doesn't show real improvement.

## Hyperparameter Tuning

Each model has parameters to fit. Increasing or decreasing them might result in overfit(in green) or underfit(in red) the model.